

All instantiations of the greedy algorithm for the shortest superstring problem are equivalent

Maksim S. Nikolaev^[0000–0003–4079–2885]

Steklov Institute of Mathematics at St. Petersburg, Russian Academy of Sciences
`makc-nicko@yandex.ru`

Abstract. In the Shortest Common Superstring problem (SCS), one needs to find the shortest superstring for a set of strings. While SCS is NP-hard and MAX-SNP-hard, the Greedy Algorithm “choose two strings with the largest overlap; merge them; repeat” achieves a constant factor approximation that is known to be at most 3.5 and conjectured to be equal to 2. The Greedy Algorithm is not deterministic, so its instantiations with different tie-breaking rules may have different approximation factors. In this paper, we show that it is not the case: all factors are equal. To prove this, we show how to transform a set of strings so that all overlaps are different whereas their ratios stay roughly the same.

Keywords: superstring · shortest common superstring · approximation · greedy algorithms · greedy conjecture

1 Introduction

In the Shortest Common Superstring problem (SCS), one is given a set of strings and needs to find the shortest string that contains each of them as a substring. Applications of this problem include genome assembly [12, 8] and data compression [3, 2, 9]. We refer the reader to the surveys [4, 7] for an overview of SCS as well as its applications and algorithms.

While SCS is NP-hard [3] and even MAX-SNP-hard [1], the Greedy Algorithm (GA) “choose two strings with the largest overlap; merge them; repeat” achieves a constant factor approximation that is proven to be less than or equal to 3.5 [6]. This factor is at least 2 (consider a dataset $\mathcal{S} = \{c(ab)^n, (ab)^n c, (ba)^n\}$) and the 30 years old *Greedy Conjecture* [9–11, 1] claims that this bound is accurate, that is, that GA is 2-approximate.

GA is not deterministic as we do not specify how to break ties in case when there are many pairs of strings with maximum overlap. For this reason, different instantiations of GA may produce different superstrings for the same input and hence they may have different approximation factors. In fact, if \mathcal{S} contains only strings of length 2 or less or if \mathcal{S} is a set of k -substrings of an unknown string, then there are instantiations of GA [5], that find *the exact* solution, whereas in general GA fails to do so.

The original Greedy Conjecture states that *any* instantiation of GA is 2-approximate. As this is still widely open, it is natural to try to prove the conjecture at least for *some* instantiations. This could potentially be easier not just

because this is a weaker statement, but also because a particular instantiation of GA may decide how to break ties by asking an almighty oracle. In this paper, we show that this weak form of Greedy Conjecture is in fact equivalent to the original one. More precisely, we show, that if *some* instantiation of GA is λ -approximate, then *all* instantiations are λ -approximate.

To prove this, we introduce the so-called *Disturbing Procedure*, that, for a given dataset $\mathcal{S} = \{s_1, \dots, s_n\}$, a parameter $m \gg n$, and a sequence of greedy non-trivial merges (merges of strings with a non-empty overlap), constructs a new dataset $\mathcal{S}' = \{s'_1, \dots, s'_n\}$, such that, for all $i \neq j$, s'_i is roughly m times longer than s_i , the overlap of s'_i and s'_j is roughly m times longer than the overlap of s_i and s_j , and the mentioned greedy sequence of non-trivial merges for \mathcal{S} is *the only* such sequence for \mathcal{S}' .

2 Preliminaries

Let $|s|$ be the length of a string s and $\text{ov}(s, t)$ be the *overlap* of strings s and t , that is, the longest string y , such that $s = xy$ and $t = yz$. In this notation, a string xyz is a *merge* of strings s and t . By ε we denote the empty string. By $\text{OPT}(\mathcal{S})$ we denote the optimal superstring for the dataset \mathcal{S} .

Without loss of generality we may assume that the set of input strings \mathcal{S} contains no string that is a substring of another. This assumption implies that in any superstring all strings occur in some order: if one string begins before another, then it also ends before. Hence, we can consider only superstrings that can be obtained from some permutation $(s_{\sigma(1)}, \dots, s_{\sigma(n)})$ of \mathcal{S} after merging adjacent strings. The length of such superstring $s(\sigma)$ is simply

$$|s(\sigma)| = \sum_{i=1}^n |s_i| - \sum_{i=1}^{n-1} |\text{ov}(s_{\sigma(i)}, s_{\sigma(i+1)})|. \quad (1)$$

Let A be an instantiation of GA (we denote this by $A \in \text{GA}$). By σ_A we denote the permutation corresponding to a superstring $A(\mathcal{S})$ constructed by A , and by $(l_A(1), r_A(1)), \dots, (l_A(n-1), r_A(n-1))$, we denote the order of merges: strings $s_{l_A(i)}$ and $s_{r_A(i)}$ are merged at step i . By the definition of GA we have

$$|\text{ov}(s_{l_A(i)}, s_{r_A(i)})| \geq |\text{ov}(s_{l_A(j)}, s_{r_A(j)})|, \quad \forall i < j < n,$$

and if, for some i , $|\text{ov}(s_{l_A(i)}, s_{r_A(i)})| = 0$, then the same holds for any $i' > i$. We denote the first such i by T_A and this is the first trivial merge (that is, one with the empty overlap), after which all the merges are trivial. Note that just before step T_A , all the remaining strings have empty overlaps, so the resulting superstring is just a concatenation of them in some order and this order does not affect the length of the result. If there were no trivial merges, we set $T_A = n$.

3 Disturbing Procedure

Here, we describe the mentioned procedure that gets rid of ties. Consider a dataset \mathcal{S} , an instantiation $A \in \text{GA}$ and a *sentinel* $\$$ — a symbol that does not



(b)

where

(2)

- Basically, we insert the string $\m before every character of s_i and then remove some $\$$'s from the beginning of the string and add some $\$$'s to its end (see Fig. 1). The purpose of this removal and addition is to disturb slightly overlaps of equal length, so there are no longer any ties in non-trivial merges.

Lemma 1. *For all $i \neq j$, $k \neq l$*

- Proof.* Let $\text{ov}(s_i, s_j)$ be $c_1 c_2 \dots c_k$. Consider the string

$$u = \$^{m-\alpha_j} c_1 \$^m \dots \$^m c_k \$^{T_A-\beta_i}.$$

Clearly, u is the overlap of s'_i and s'_j and $|u| = (m+1)k - \alpha_j + T_A - \beta_i$. Also, if $|\text{ov}(s_i, s_j)| = 0$ then $\text{ov}(s'_i, s'_j) = \$^{\min\{T_A - \beta_i, m - \alpha_j\}}$.

To prove the last statement, note that $\alpha_j + \beta_i \leq 2T_A < m$ and

$$|\text{ov}(s'_i, s'_j)| > (m+1)|\text{ov}(s_k, s_l)| + T_A \geq |\text{ov}(s'_k, s'_l)|.$$

Lemma 2. *Let $B \in \text{GA}$. Then $T_A = T'_A = T'_B$ and the first $T_A - 1$ merges are the same for both instantiations.*

Proof. We prove by induction that $l_A(t) = l'_A(t) = l'_B(t)$ and $r_A(t) = r'_A(t) = r'_B(t)$ for all $t < T_A$.

Case $t = 1$. As A is greedy, then $k_1 := |\text{ov}(s_{l_A(1)}, s_{r_A(1)})| \geq |\text{ov}(s_i, s_j)|$, for all $i \neq j$, $(i, j) \neq (l_A(1), r_A(1))$. Hence

$$\begin{aligned} |\text{ov}(s'_i, s'_j)| &\leq (m+1)k_1 - \alpha_j + T_A - \beta_i < \\ &< (m+1)k_1 - 1 + T_A - 1 = |\text{ov}(s'_{l_A(1)}, s'_{r_A(1)})|, \end{aligned}$$

and $l'_A(1) = l'_B(1) = l_A(1)$ as well as $r'_A(1) = r'_B(1) = r_A(1)$.

Suppose that the statement holds for all $t \leq t' < T_A - 1$. Note that at moment $t = t' + 1$ the sum $\alpha_j + \beta_i$ is strictly greater than $2t$ unless $(i, j) = (l_A(t), r_A(t))$. Similarly to the base case, we have

$$\begin{aligned} |\text{ov}(s'_i, s'_j)| &\leq (m+1)k_t - \alpha_j + T_A - \beta_i < \\ &< (m+1)k_t - t + T_A - t = |\text{ov}(s'_{l_A(t)}, s'_{r_A(t)})|, \end{aligned}$$

where $k_t = |\text{ov}(s_{l_A(t)}, s_{r_A(t)})|$, and the induction step is proven.

Now note that starting from step T_A all the remaining strings in \mathcal{S} have empty overlaps and hence so do the remaining strings in \mathcal{S}' , as for all of them $\beta_i = T_A$ and the minimum in Lemma 1.2 is equal to zero. Thus, $T_A = T'_A = T'_B$ and the lemma is proven.

Corollary 1. *As all non-trivial merges coincide, $|A(\mathcal{S}')| = |B(\mathcal{S}')|$.*

4 Equivalence of instantiations

Theorem 1. *If some instantiation A of GA achieves λ -approximation, then so does any other instantiation.*

Proof. Assume the opposite and consider $B \in \text{GA}$ as well as a dataset \mathcal{S} such that $|B(\mathcal{S})| > \lambda|\text{OPT}(\mathcal{S})|$. Let $\mathcal{S}' = \mathcal{S}'(B, m)$ be the corresponding disturbed dataset, where $m > 2n$ will be specified later.

Note that $|s'_i|/m \rightarrow |s_i|$ and $|\text{ov}(s'_i, s'_j)|/m \rightarrow |\text{ov}(s_i, s_j)|$ as m approaches infinity, thanks to Lemma 1.1–2. Then $|\text{OPT}(\mathcal{S}')|/m \rightarrow |\text{OPT}(\mathcal{S})|$, since

$$\begin{aligned} \frac{1}{m}|\text{OPT}(\mathcal{S}')| &= \frac{1}{m} \min_{\sigma} \left\{ \sum_{i=1}^n |s'_i| - \sum_{i=1}^{n-1} |\text{ov}(s'_{\sigma(i)}, s'_{\sigma(i+1)})| \right\} \rightarrow \\ &\rightarrow \min_{\sigma} \left\{ \sum_{i=1}^n |s_i| - \sum_{i=1}^{n-1} |\text{ov}(s_{\sigma(i)}, s_{\sigma(i+1)})| \right\} = |\text{OPT}(\mathcal{S})|, \end{aligned}$$

$|B(\mathcal{S}')|/m \rightarrow |B(\mathcal{S})|$ and hence $|A(\mathcal{S}')|/m \rightarrow |B(\mathcal{S})|$, by Corollary 1.

As $|B(\mathcal{S})| - \lambda|\text{OPT}(\mathcal{S})| > 0$, we can choose m so that $|B(\mathcal{S}')| - \lambda|\text{OPT}(\mathcal{S}')|$ as well as $|A(\mathcal{S}')| - \lambda|\text{OPT}(\mathcal{S}')|$ are positive. Hence A is not λ -approximate.

Corollary 2. *To prove (or disprove) the Greedy Conjecture, it is sufficient to consider datasets satisfying one of the following three properties:*

1. *there are no ties between non-empty overlaps, that is, datasets where all the instantiations of the greedy algorithm work the same;*
2. *there are no empty overlaps: $\text{ov}(s_i, s_j) \neq \varepsilon, \forall i \neq j$;*
3. *all non-empty overlaps are (pairwise) different: $|\text{ov}(s_i, s_j)| \neq |\text{ov}(s_k, s_l)|$, for all $i \neq j, k \neq l, (i, j) \neq (k, l)$.*

Proof. 1. Follows directly from the proof of Theorem 1, as we always can use the dataset \mathcal{S}' instead of \mathcal{S} .

2. Append $\$$ to each string of \mathcal{S}' . Then, every two strings have non-empty overlap that at least contains $\$$, and in general $T_A = T'_A = T'_B$ from Lemma 2 does not hold (T'_A and T'_B are always n). However, the first T_A merges are still the same and after them all the remaining strings have overlaps of length 1 and then the lengths of the final solutions are the same as well.
3. Append $\$^{n(T_A - \beta_i)}$ to each string of \mathcal{S}' instead of $\$^{T_A - \beta_i}$. Then

$$|\text{ov}(s'_i, s'_j)| = (m + 1)|\text{ov}(s_i, s_j)| - \alpha_j + nT_A - n\beta_i,$$

provided m is large enough, and $\alpha_j + n\beta_i \neq \alpha_k + n\beta_l$ if $(i, j) \neq (k, l)$. Repeating the proofs of Lemmas 1 and 2 with this version of \mathcal{S}' , we obtain this statement of the lemma.

5 Conclusion

In this paper we revealed the equivalence of greedy algorithms for the shortest common superstring problem. This means, in particular, that proving or disproving the Greedy Conjecture is difficult not due to non-deterministic nature of the Greedy Algorithm, but due to the complexity of the overlaps structure.

Acknowledgments

Many thanks to Alexander Kulikov for valuable discussions and proofreading the text.

References

1. Avrim Blum, Tao Jiang, Ming Li, John Tromp, and Mihalis Yannakakis. Linear approximation of shortest superstrings. In *STOC 1991*, pages 328–336. ACM, 1991.
2. John Gallant. *String compression algorithms*. PhD thesis, Princeton, 1982.

3. John Gallant, David Maier, and James A. Storer. On finding minimal length superstrings. *J. Comput. Syst. Sci.*, 20(1):50–58, 1980.
4. Theodoros P. Gevezes and Leonidas S. Pitsoulis. *The shortest superstring problem*, pages 189–227. Springer, 2014.
5. Alexander Golovnev, Alexander S Kulikov, Alexander Logunov, Ivan Mihajlin, and Maksim Nikolaev. Collapsing superstring conjecture. In *Approximation, Randomization, and Combinatorial Optimization. Algorithms and Techniques (APPROX/RANDOM 2019)*. Schloss Dagstuhl-Leibniz-Zentrum fuer Informatik, 2019.
6. Haim Kaplan and Nira Shafrir. The greedy algorithm for shortest superstrings. *Inf. Process. Lett.*, 93(1):13–17, 2005.
7. Marcin Mucha. A tutorial on shortest superstring approximation. <https://www.mimuw.edu.pl/~much/teaching/aa2008/ss.pdf>. Online; accessed 24 May 2021.
8. Pavel A. Pevzner, Haixu Tang, and Michael S. Waterman. An eulerian path approach to DNA fragment assembly. *Proc. Natl. Acad. Sci. U.S.A.*, 98(17):9748–9753, 2001.
9. James A. Storer. *Data compression: methods and theory*. Computer Science Press, Inc., 1987.
10. Jorma Tarhio and Esko Ukkonen. A greedy approximation algorithm for constructing shortest common superstrings. *Theor. Comput. Sci.*, 57(1):131–145, 1988.
11. Jonathan S. Turner. Approximation algorithms for the shortest common superstring problem. *Inf. Comput.*, 83(1):1–20, 1989.
12. Michael S. Waterman. *Introduction to computational biology: maps, sequences and genomes*. CRC Press, 1995.